Exploring Parallel MCTS on Chess Game

曾正豪 0716325
CS NYCU
Hsinchu, Taiwan

張宸愷 0710018
EECSHP NYCU
Hsinchu, Taiwan

Figure 1: Illustration for a single step of MCTS

ABSTRACT

A project proposal for the course 'Parallel Programming Fall 2021'. We decided to explore the parallelization of Monte Carlo Tree Search using the techniques and knowledge we have learned in this course. We will use quantitative benchmarks to compare different approaches to solve this kind of parallelization.

KEYWORDS

MCTS, parallel programming, Pthreads, CUDA, Chess

1 INTRODUCTION

'Monte Carlo Tree Search', abbreviated in this text from now on as MCTS, is a probability sampling based tree search method for many applications. One of the most famous application of MCTS is AlphaGO [5]. It uses MCTS with 2 other neural networks to play Go. An early version of AlphaGo was tested on hardware with various numbers of CPUs and GPUs, running in asynchronous or distributed mode. It was tested with search threads from 12 to 64, number of CPUs from 48 to 1920, and number of GPUs from 1 to 280.

And in 2016, it changed to use TPUs (tensor processing units) as its computing unit. In recent years, it keeps beating many go players. Overall, MCTS is an algorithm that can be highly parallelized because of the high number of simulations. Hence, we decided to use MCTS as the topic of our final project.

2 STATEMENT OF PROBLEM

One of our teammates had taken the course AI capstone, and during that course he had been doing a final project about playing a 3D version of connect-4 using upper confidence bound MCTS, i.e. UCB-MCTS. He noticed that when he uses root parallelization, the performance of the MCTS decreases dramatically. The multi-threaded version with 8 threads didn't even manage to beat the single-threaded one. The teammate's guess is that it's because of false sharing. Thus, we want to explore further on ways to improve multithreaded performance, and hopefully use this new improved MCTS to play on well-known games such as Chess, instead of some obscure game that our teammate played back then.

3 PROPOSED APPROACHES

3.1 Overall design

Figure 2 is our system design. Component descriptions:

- Chess game module: Deal with the game state of Chess, once it receives a legal move, it will change its game state and return it.
- Arbiter and round-robin tournament maker: It will coordinate the game player and the game module. It checks if one move is legal and whether the game ends.
- 3 MCTS versions: To calculate the best move next. It will receive a game state and return its best move. The difference between these 3 versions is their parallel approach.
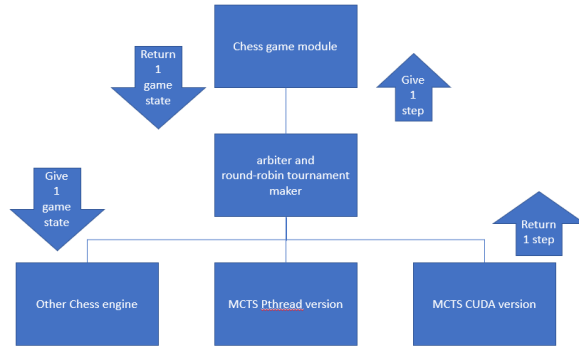
Figure 2: Our system diagram

- Additional: Maybe we will compete with other Chess engines to see the performance pitted against other state of the art, if we have enough time.

The performance of each method will depend on the games won and the total amount of expanded nodes.

3.2 parallelization details

There are four ways to parallelize traditional UCB-MCTS mentioned in [1], and illustrated in figure 3. Leaf parallelization, root parallelization, tree parallelization with global mutex, and tree parallelization with local mutexes.
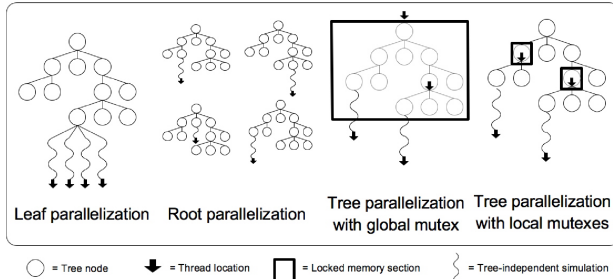

Figure 3: Ways of parallelizing MCTS

We will do comparison between GPU-based leaf parallelization, CPU-based root parallelization, P-UCT (from [2]).

4 LANGUAGE SELECTION

We will use C++ as our main programming language, together with Pthreads and CUDA, since we want to compare the performance of CPU and GPU on parallel programs.

5 RELATED WORK

[4] is a introductory view on UCB-MCTS with additional enhancements e.g. rapid action value estimate (RAVE), all moves as first (AMAF), etc. A pretty detailed work [3] regarding current developments on different MCTS, including some exotic methods that uses virtual loss [2] (i.e. do simulation before the previous backprop result is known). Different methods for enhancing the capability of

MCTS on board games were explored in [6]. [1][6] also outlined some of the parallelization techniques for MCTS, root parallelization, leaf parallelization, and tree parallelization, etc. [6] also outlined a way for a modularized MCTS algorithm that can scale up as the number of computer nodes grow.

6 EXPECTED RESULTS

When using tree parallelization, there will be two methods. One is we use a global mutex to protect all the nodes in the tree. The other is to give each node a mutex. The former will be much slower because many threads are contending for a single lock, while the latter will use significantly more memory because each node has a lock, but the performance will be better. The result of the proposed solution from [2] will probably beat root parallelization and leaf parallelization.

The MCTS will probably benefit the most from the high throughput of the GPU, because it requires an enormous amount of simulation (rollouts) to converge to the best move.

7 TIMETABLE

Below are the rough due dates for the individual components of our project.

- 10/27 proposal
- 11/05 chess engine wrapper and game module
- 11/20 Pthread version
- 12/10 CUDA version
- 12/17 Performance analysis
- 12/31 Finalize report and source code

8 APPENDICES

REFERENCES

[1] Guillaume M.J-B. Chaslot, Mark H.M. Winands, and H. Jaap van den Herik. 2008. Parallel Monte-Carlo Tree Search.
[2] Anji Liu, Jianshu Chen, Mingze Yu, Yu Zhai, Xuewen Zhou, and Ji Liu. 2020. Watch the Unobserved: A Simple Approach to Parallelizing Monte Carlo Tree Search. In *International Conference on Learning Representations.* https://openreview.net/forum?id=BJlQtJSKDB
[3] Anji Liu, Yitao Liang, Ji Liu, Guy Van den Broeck, and Jianshu Chen. 2020. On Effective Parallelization of Monte Carlo Tree Search. arXiv:2006.08785 [cs.LG]
[4] Magnuson and Max. 2015. Monte Carlo Tree Search and Its Applications. *University of Minnesota, Morris Undergraduate Journal* 2, 2 (2015).
[5] David Silver, Aja Huang, Christopher Maddison, Arthur Guez, Laurent Sifre, George Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. 2016. Mastering the game of Go with deep neural networks and tree search. *Nature* 529 (01 2016), 484–489. https://doi.org/10.1038/nature16961
[6] Martin Weigel. 2017. Monte Carlo methods for massively parallel computers. arXiv:1709.04394 [physics.comp-ph]

A ONLINE RESOURCES

(1) Monte Carlo tree search (https://en.wikipedia.org/wiki/Monte_Carlo_tree_search)